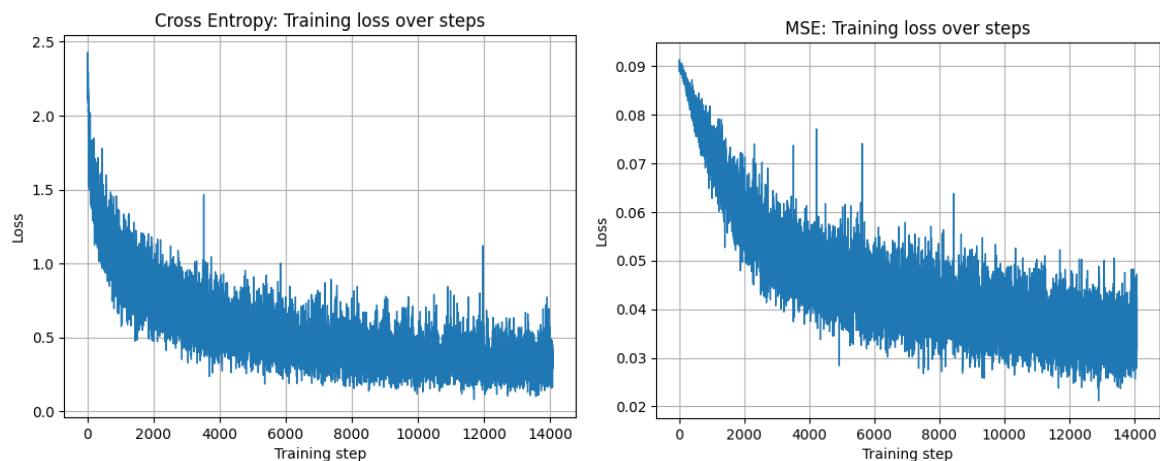


Project 2

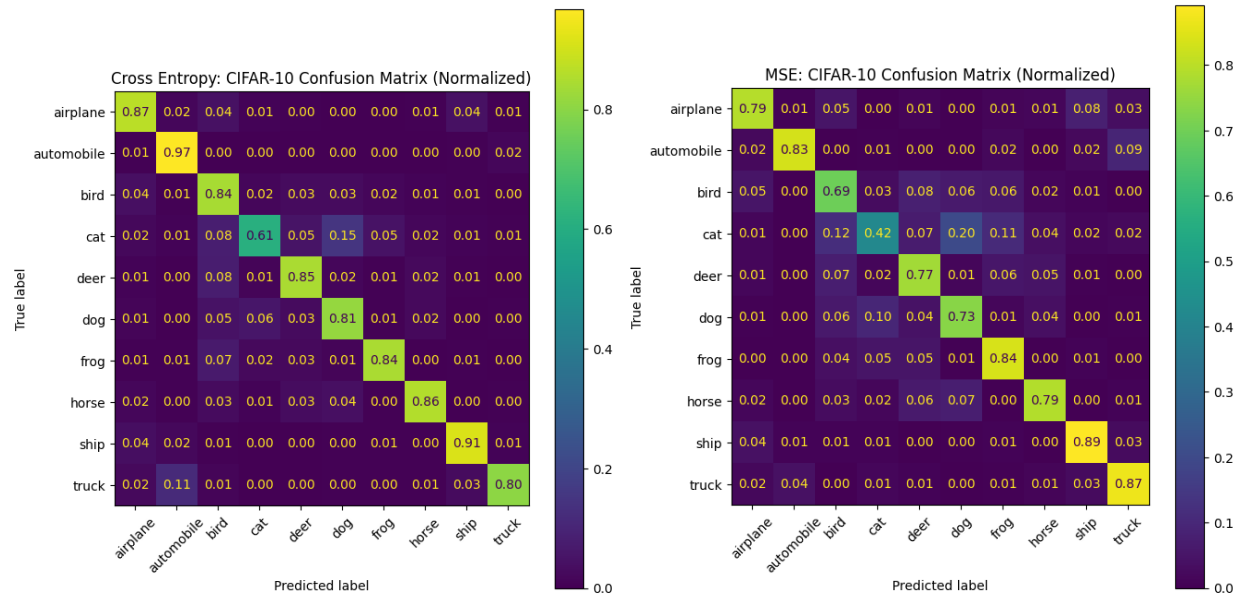
Machine Perception

John Brooks

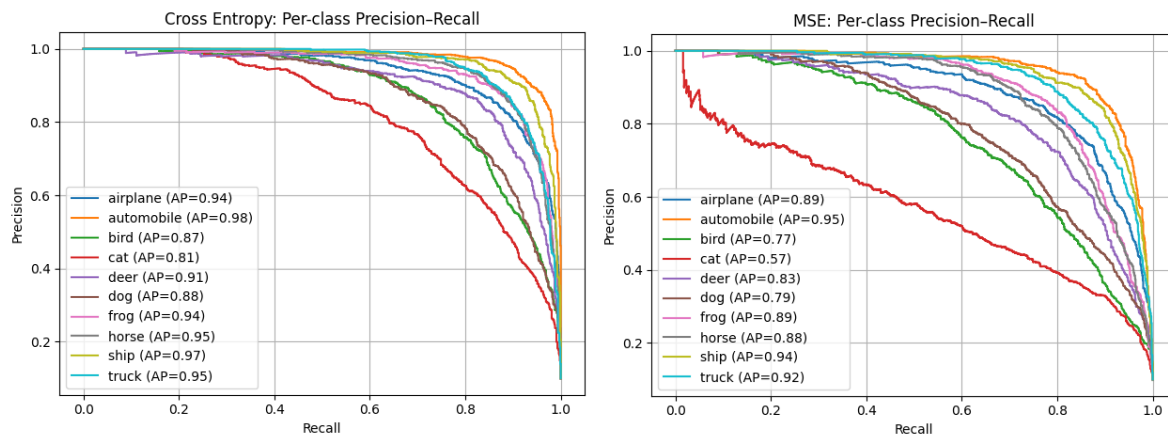
For my implementation of the AlexNet architecture to classify to the CIFAR-10 dataset I developed two implementations. The first implementation utilizing Cross Entropy Loss and the second utilizing mean squared error loss. They both utilize stochastic gradient descent as the optimizer. I stuck with a batch size 64 and only trained each run to 20 epochs. Both models baseline setup utilized a learning rate of 0.005, a weight decay of 0.005, and a momentum of 0.9. The cross-entropy model achieved an accuracy of 82.4% on the training data and an 82.12% on the test data. The mean squared error model achieved an accuracy of 75.42% on the training data and a 75.15% on the test data. The training loss as a function of steps is graphed as such for cross-entropy and MSE:



The confusion matrix is shown on the next page for both models:



Shown below is the precision recall curve:



As can be seen in the MSE precision recall curve, cat had a very different precision compared to the rest of the classes in the model. Shown on the next page we can see a picture of a puppy that was classified as a dog by both models as well as the top 5 possible predictions the model thought the subject of the image was:

Top-1 prediction: dog (0.959)

Top-5 predictions:

dog	0.959
bird	0.022
frog	0.009
deer	0.005
cat	0.005

Top-1 prediction: dog (0.596)

Top-5 predictions:

dog	0.596
cat	0.149
horse	0.126
bird	0.072
truck	0.019

(Cross Entropy) Predicted: dog



(MSE) Predicted: dog



As can be seen in the above images, the cross-entropy model is far more confident that the dog is a dog compared to the MSE model. I also experimented with testing the models on planes which neither classified correctly and tested the models on a pickup truck (which does not exist in the dataset) and both models identified the vehicle as an automobile.